

Postmeet.

The last mile of meetings.

Every meeting produces commitments. Most of them evaporate. Postmeet turns a raw transcript into distributed, owned action — in about five seconds, with zero setup.

Shubham Mittal · AI Product Manager



Top 10 Finalist → 2nd Runner-Up

PromptWars Chennai 2026 · Google for Developers × Hack2Skill

Live demo: postmeet.onrender.com

Repo: github.com/Shubham-33/postmeet

— SEE IT WORK

Paste a transcript → the board writes itself.

SUMMARY

Q3 planning meeting with discussions on auth refactor, dark mode, and design-token branch.

Email MOM to team →

Copy MOM as text

All Calendars →

All Emails →

Decisions01

We decided to ship dark mode in the v3 release.

JA Jamie02

Merge the design-token branch

it's blocking Tara's work

Tue, Jul 14, 2026

jamie@example.com

Calendar →

Email →

Tag Tara

it's blocking Tara's work

+ add date jamie@example.com

Calendar →

Email →

MA Marcus02

Have the PR ready for reviews

for Q3 planning

Fri, Jul 10, 2026

marcus@example.com

Calendar →

Email →

Write the migration notes

so QA isn't blocked

Mon, Jul 13, 2026

marcus@example.com

Calendar →

Email →

Real output from the live app — extracted in ~3s: summary, decisions, and one column per owner, each card one click from a prefilled Calendar event or Gmail draft.

Paste a meeting → get owned action → distribute in one click.

INPUT

Drop it in

Paste a transcript, a public Google Doc URL, or a file. No account.

EXTRACT

AI structures it

Summary, decisions, and action items — each with owner, email, due date, context.

DISTRIBUTE

One-click send

Each item opens a prefilled Google Calendar event or Gmail draft. Just Save / Send.

The “no-OAuth trick”: prefilled Calendar / Gmail URLs open in the user's already-logged-in Google. **Time-to-value ~10 seconds — no integration, no admin approval, no scope review.**

Meetings manufacture commitments. Then the commitments go missing.

~70%

of meeting action items are never completed — they live in someone's notes, a Slack thread, a doc nobody reopens.

who

Knowledge workers spend a huge share of the week in meetings — what matters is the follow-through, not the recording.

gap

The pain isn't capturing what was said. It's doing what was decided.

cost

Dropped commitments = re-work, missed deadlines, and eroded trust.

Everyone is fighting over capture. The last mile is wide open.

CROWDED — “CAPTURE”

Transcribe & summarize

Otter, Fireflies, Fathom, Granola — all competing on better transcripts. Table stakes now. The commitments they surface still die in a doc.

OPEN — “THE LAST MILE”

Distribute & act

Getting each commitment into the tool the owner actually lives in — their calendar, their inbox — is where follow-through is won or lost.

The blocker on the last mile isn't a smarter model — it's **setup friction**. So the wedge is zero-setup distribution.

Reliable structure out of an unreliable medium.

transcript → **LLM · JSON mode** → schema-validated JSON → board + prefilled links

Structured output

A fixed JSON contract (summary / decisions / action items), pinned in the system prompt and decoded by a tolerant parser. One call, no retry loop.

Guardrails

Extracts only explicit commitments — no invented tasks. Empty fields when data is absent. The difference between useful and hallucinated.

Graceful degradation

A model fallback chain: if the primary is cold or errors, it drops to a backup so the request still succeeds.

I shipped the smaller model on purpose.

REJECTED AS PRIMARY

Llama 3.1 70B

latency	9s-46s, inconsistent
reliability	cold-starts, timeouts
quality	marginally higher

SHIPPED AS PRIMARY

Llama 3.1 8B

latency	~2.5s, consistent
reliability	steady across runs
quality	comparable

Model quality is an **input** to product value, not the goal. A demo that always works in two seconds beats a “smarter” one that sometimes hangs — reliability builds trust faster than marginal accuracy.

One north star, a funnel beneath it, guardrails around it.

NORTH-STAR METRIC

Commitment Completion Rate

Of the action items surfaced, what share actually get done? Everything else is a proxy.

INPUT METRICS – THE FUNNEL

- Activation: % sessions that extract → distribute ≥ 1 item
- Time-to-value: seconds paste → first send
- Extraction quality: precision / recall vs. eval set
- Retention: repeat use / week

GUARDRAIL METRICS

- Hallucination rate (invented items)
- Edit rate (proxy for accuracy)
- Error / failure rate

Vibes don't scale. Evals do.

- set** Build a golden eval set: labelled transcripts with the “true” decisions and action items.
- score** Measure precision (did we invent tasks?), recall (did we miss commitments?), and field accuracy.
- use** That's how I'd choose 8B vs 70B, catch regressions, and set a launch quality bar — not by gut.

THE COMPOUNDING LOOP

Every edit is a label

Users correct the extracted items before sending. Those corrections are free, in-domain training data → fine-tune on accepted output → org-specific accuracy competitors can't match without the same data. The eval loop becomes the moat.

What I deliberately left out — and why that's the point.

MVP scope is a hypothesis test, not a feature checklist. I cut everything not needed to validate one thing: will people extract, then actually distribute?

CUT · PERSISTENCE

No database. A refresh clears everything — and “no data stored” removes privacy friction for testers.

CUT · REAL INTEGRATIONS

The URL-prefill trick proves value without OAuth. Real APIs are a scaling investment, made once the loop is proven.

CUT · ACCOUNTS & TEAMS

Single-session, no login. Multi-user is a growth concern, not a validation one.

KEPT · THE CORE LOOP + TRUST

Extraction reliability, one-click distribution, honest guardrails, 100%-tested logic.

From a one-click helper to the system of record for follow-through.

NOW · VALIDATE

The core loop

- Paste → extract → distribute
- Zero setup, no OAuth
- Reliable structured output

NEXT · AUTO-CAPTURE

Capture everywhere

- Extension + AI notetaker
- Auto-joins Zoom / Teams / Meet
- Native Gmail, Outlook, Calendar

LATER · OWN THE OUTCOME

Proactive follow-through

- Persistence + team workspaces
- Nudges: “did you ship X?”
- Completion tracking → closes the loop

Moat = the feedback loop. Accepted & edited action items become proprietary training data → accuracy that compounds with usage.

I don't just use AI tools. I make product decisions about them.

Problem framing

Found the non-obvious wedge — the last mile, not capture; friction, not model size.

AI judgment

Model tradeoffs, evals, guardrails, graceful degradation — decisions, with reasons.

Metrics rigor

A north star, an input funnel, and guardrails — measurement designed in.

Prioritization

Scoped an MVP as a hypothesis test; cut with intent; shipped what de-risks it.

Vision

A credible path from a one-click tool to a defensible product with a data moat.

End-to-end

Built, tested (100% coverage), and deployed it solo — insight to live URL.